

ASSIGNMENT - 4

Github link: <https://github.com/tanikonda7441/Assignment4>

1. Pandas

Importing all the required libraries to work with tabular data and also implement algorithms.

```
In [9]: #importing the required Libraries to work with Tabular data and also to implement algorithms
import warnings
import numpy as np
import pandas as pd
import seaborn as sns
from sklearn import preprocessing
import matplotlib.pyplot as plt
from scipy.stats.stats import pearsonr
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, recall_score, precision_score, classification_report, confusion_matrix
warnings.filterwarnings("ignore")
```

Question: 1

1. Read the provided CSV file 'data.csv'. <https://drive.google.com/drive/folders/1h8C3mLsso-R-sI0LsvoYwPlzy2fJ4IOF?usp=sharing>
2. Show the basic statistical description about the data.
3. Check if the data has null values. a. Replace the null values with the mean
4. Select at least two columns and aggregate the data using: min, max, count, mean.
5. Filter the dataframe to select the rows with calories values between 500 and 1000.
6. Filter the dataframe to select the rows with calories values > 500 and pulse < 100.
7. Create a new "df_modified" dataframe that contains all the columns from df except for "Maxpulse".
8. Delete the "Maxpulse" column from the main df dataframe
9. Convert the datatype of Calories column to int datatype.
10. Using pandas create a scatter plot for the two columns (Duration and Calories).

```
In [7]: #1. Read the provided CSV file 'data.csv'. https://drive.google.com/drive/folders/1h8C3mLsso-R-sI0LsvoYwPlzy2fJ4IOF?usp=sharing
df = pd.read_csv("data.csv")
df.head()
```

using the Pandas library to read a CSV file named "data.csv" and store its contents in a DataFrame object called df.

The `pd.read_csv()` function is a method provided by the Pandas library that reads the CSV file and creates a DataFrame object from it.

The `df.head()` function is then called to display the first five rows of the DataFrame.

In [7]: #1. Read the provided csv file 'data.csv'. <https://drive.google.com/drive/folders/1h8C3mLss0-R-sIOLsvoYwPLzy2fJ4IOF?usp=sharing>

```
df = pd.read_csv("data.csv")
df.head()
```

Out[7]:

	Duration	Pulse	Maxpulse	Calories
0	60	110	130	409.1
1	60	117	145	479.0
2	60	103	135	340.0
3	45	109	175	282.4
4	45	117	148	406.0

In [8]: #2. Show the basic statistical description about the data.

```
df.describe()
```

Out[8]:

	Duration	Pulse	Maxpulse	Calories
count	169.000000	169.000000	169.000000	164.000000
mean	63.846154	107.461538	134.047337	375.790244
std	42.299949	14.510259	16.450434	266.379919
min	15.000000	80.000000	100.000000	50.300000
25%	45.000000	100.000000	124.000000	250.925000
50%	60.000000	105.000000	131.000000	318.600000
75%	60.000000	111.000000	141.000000	387.600000
max	300.000000	159.000000	184.000000	1860.400000

The `df.describe()` method is a built-in function in Pandas that generates descriptive statistics of the DataFrame `df`. It includes the count, mean, standard deviation, minimum, 25th percentile, 50th percentile (median), 75th percentile, and maximum values for each numeric column in the dataframe.

jupyter Machine_learning_Assignment-4 (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3 (ipykernel) ○

In [10]: #3. Check if the data has null values.

```
df.isnull().any()
```

Out[10]: Duration False
Pulse False
Maxpulse False
Calories True
dtype: bool

In [11]: Replace the null values with the mean

```
df.fillna(df.mean(), inplace=True)  
df.isnull().any()
```

Out[11]: Duration False
Pulse False
Maxpulse False
Calories False
dtype: bool

In [12]: #4. Select at Least two columns and aggregate the data using: min, max, count, mean.

```
df.agg({'Maxpulse':['min','max','count','mean'],'Calories':['min','max','count','mean']})
```

Out[12]:

	Maxpulse	Calories
min	100.000000	50.300000
max	184.000000	1860.400000
count	169.000000	169.000000
mean	134.047337	375.790244

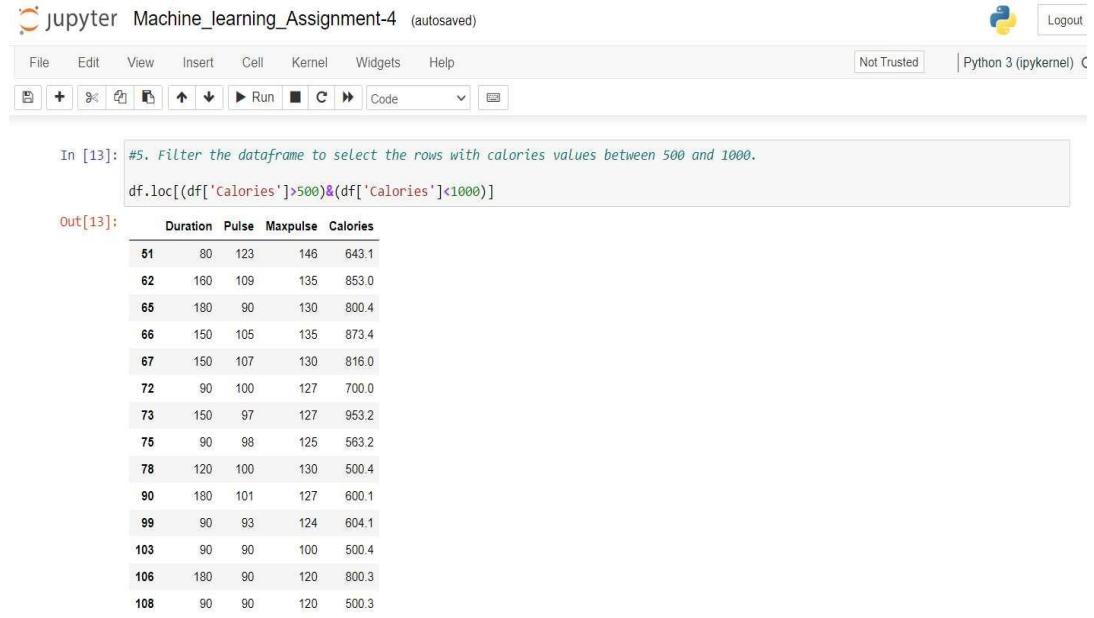
The `df.isnull().any()` code checks whether there are any missing values (also known as `NaN` or `null` values) in each column of the DataFrame `df`.

The first line of this code fills in any missing values in the DataFrame df with the mean value of each column using the `fillna()` method.

The second line of this code uses `df.isnull().any()` to check if there are any missing values remaining in the DataFrame after filling in the

missing values with the mean.

This code uses the agg() method to perform aggregate calculations on the Maxpulse and Calories columns of the DataFrame df.



In [13]: #5. Filter the dataframe to select the rows with calories values between 500 and 1000.
df.loc[(df['Calories']>500)&(df['Calories']<1000)]

Out[13]:

	Duration	Pulse	Maxpulse	Calories
51	80	123	146	643.1
62	160	109	135	853.0
65	180	90	130	800.4
66	150	105	135	873.4
67	150	107	130	816.0
72	90	100	127	700.0
73	150	97	127	953.2
75	90	98	125	563.2
78	120	100	130	500.4
90	180	101	127	600.1
99	90	93	124	604.1
103	90	90	100	500.4
106	180	90	120	800.3
108	90	90	120	500.3

This code uses the loc[] method of the DataFrame df to select rows where the value in the Calories column is greater than 500 and less than 1000.

In [14]: #6. Filter the dataframe to select the rows with calories values > 500 and pulse < 100.
df.loc[(df['Calories']>500)&(df['Pulse']<100)]

Out[14]:

	Duration	Pulse	Maxpulse	Calories
65	180	90	130	800.4
70	150	97	129	1115.0
73	150	97	127	953.2
75	90	98	125	563.2
99	90	93	124	604.1
103	90	90	100	500.4
106	180	90	120	800.3
108	90	90	120	500.3

In [15]: #7. Create a new "df_modified" dataframe that contains all the columns from df except for "Maxpulse".
df_modified = df[['Duration','Pulse','Calories']]
df_modified.head()

Out[15]:

	Duration	Pulse	Calories
0	80	110	409.1
1	60	117	479.0
2	60	103	340.0
3	45	109	282.4
4	45	117	406.0

This code uses the loc[] method of the DataFrame df to select rows where the value in the Calories column is greater than 500 and the value in the Pulse column is less than 100.

This code creates a new DataFrame `df_modified` that includes only the Duration, Pulse, and Calories columns of the original DataFrame `df`.

The `head()` method is then called on the `df_modified` DataFrame to display the first five rows of the new DataFrame.

```
In [16]: #8. Delete the "Maxpulse" column from the main df dataframe  
del df['Maxpulse']
```

```
In [17]: df.head()
```

```
Out[17]:   Duration  Pulse  Calories  
0          60     110    409.1  
1          60     117    479.0  
2          60     103    340.0  
3          45     109    282.4  
4          45     117    406.0
```

```
In [18]: df.dtypes
```

```
Out[18]: Duration    int64  
Pulse      int64  
Calories   float64  
dtype: object
```

```
In [19]: #9. Convert the datatype of Calories column to int datatype.
```

```
df['Calories'] = df['Calories'].astype(np.int64)  
df.dtypes
```

```
Out[19]: Duration    int64  
Pulse      int64  
Calories   int64  
dtype: object
```

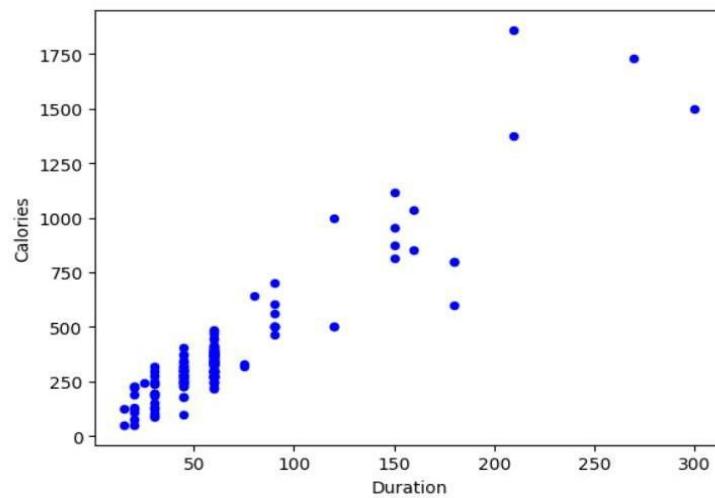
After running this code, the Maxpulse column will no longer be present in the DataFrame `df`. If the column was present in the original DataFrame `df`, then it has now been removed permanently. The resulting DataFrame will have one less column than the original DataFrame.

The `astype()` method is a Pandas method that is used to cast a column of a DataFrame to a specified data type. In this case, the Calories column is being cast to the 64-bit integer data type using the `np.int64` NumPy data type.

The `dtypes` attribute of the DataFrame is then used to display the data types of each column in the DataFrame. After running this code, the Calories column will have a data type of `int64`.

```
In [20]: #10. Using pandas create a scatter plot for the two columns (Duration and Calories).
df.plot.scatter(x='Duration',y='Calories',c='blue')

Out[20]: <Axes: xlabel='Duration', ylabel='Calories'>
```



a scatter plot of the Duration and Calories columns in the DataFrame df will be displayed, where each data point is represented by a blue marker. The x-axis will correspond to the Duration column values and the y-axis to the Calories column values.

Question: 2

Titanic Dataset

1. Find the correlation between 'survived' (target column) and 'sex' column for the Titanic use case in class. a. Do you think we should keep this feature?
2. Do at least two visualizations to describe or show correlations.
3. Implement Naïve Bayes method using scikit-learn library and report the accuracy

```
In [28]: #Loading the data file into te program
df=pd.read_csv("train.csv")

df.head()
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th... Heikkinen, Miss. Laina	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Allen, Mr. William Henry	male	35.0	1	0	113803	53.1000	C123	S
4	5	0	3			35.0	0	0	373450	8.0500	NaN	S

```
In [29]: #converted categorical data to numerical values for correlation calculation
label_encoder = preprocessing.LabelEncoder()
df['Sex'] = label_encoder.fit_transform(df.Sex.values)

#Calculation of correlation for 'Survived' and 'Sex' in data
correlation_value= df['Survived'].corr(df['Sex'])

print(correlation_value)
-0.5433513806577547
```

Ans: Yes, we should keep the 'Survived' and 'Sex' features helps classify the data accurately

```
In [30]: #print correlation matrix
matrix = df.corr()
print(matrix)
```

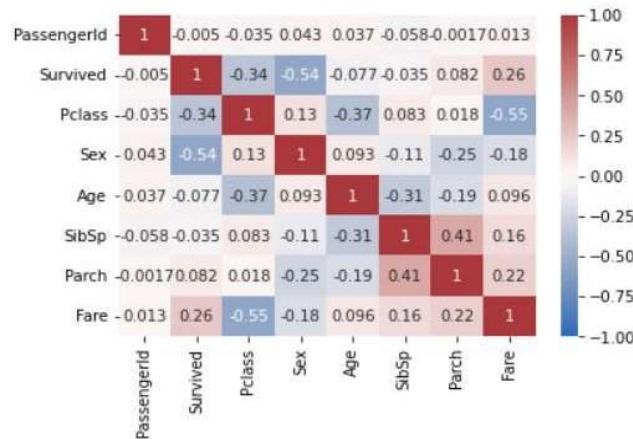
	PassengerId	Survived	Pclass	Sex	Age	SibSp	\
PassengerId	1.000000	-0.005007	-0.035144	0.042939	0.036847	-0.057527	
Survived	-0.005007	1.000000	-0.338481	-0.543351	-0.077221	-0.035322	
Pclass	-0.035144	-0.338481	1.000000	0.131900	-0.369226	0.083081	
Sex	0.042939	-0.543351	0.131900	1.000000	0.093254	-0.114631	
Age	0.036847	-0.077221	-0.369226	0.093254	1.000000	-0.308247	
SibSp	-0.057527	-0.035322	0.083081	-0.114631	-0.308247	1.000000	
Parch	-0.001652	0.081629	0.018443	-0.245489	-0.189119	0.414838	
Fare	0.012658	0.257307	-0.549500	-0.182333	0.096067	0.159651	
		Parch	Fare				
PassengerId	-0.001652	0.012658					
Survived	0.081629	0.257307					
Pclass	0.018443	-0.549500					
Sex	-0.245489	-0.182333					
Age	-0.189119	0.096067					
SibSp	0.414838	0.159651					
Parch	1.000000	0.216225					
Fare	0.216225	1.000000					

```
In [31]: # One way of visualizing correlation matrix in form of spread chart
df.corr().style.background_gradient(cmap="Reds")
```

Out[31]:

	PassengerId	Survived	Pclass	Sex	Age	SibSp	Parch	Fare
PassengerId	1.000000	-0.005007	-0.035144	0.042939	0.036847	-0.057527	-0.001652	0.012658
Survived	-0.005007	1.000000	-0.338481	-0.543351	-0.077221	-0.035322	0.081629	0.257307
Pclass	-0.035144	-0.338481	1.000000	0.131900	-0.369226	0.083081	0.018443	-0.549500
Sex	0.042939	-0.543351	0.131900	1.000000	0.093254	-0.114631	-0.245489	-0.182333
Age	0.036847	-0.077221	-0.369226	0.093254	1.000000	-0.308247	-0.189119	0.096067
SibSp	-0.057527	-0.035322	0.083081	-0.114631	-0.308247	1.000000	0.414838	0.159651
Parch	-0.001652	0.081629	0.018443	-0.245489	-0.189119	0.414838	1.000000	0.216225
Fare	0.012658	0.257307	-0.549500	-0.182333	0.096067	0.159651	0.216225	1.000000

```
In [32]: #Second form of visualizing correlation matrix using heatmap() from seaborn
sns.heatmap(matrix, annot=True, vmax=1, vmin=-1, center=0, cmap='vlag')
plt.show()
```



```
In [33]: #Loaded data files test and train and merged files
train_raw = pd.read_csv('train.csv')
test_raw = pd.read_csv('test.csv')
train_raw['train'] = 1
test_raw['train'] = 0
df = train_raw.append(test_raw, sort=False)
features = ['Age', 'Embarked', 'Fare', 'Parch', 'Pclass', 'Sex', 'SibSp']
target = 'Survived'
df = df[features + [target] + ['train']]
df['Sex'] = df['Sex'].replace(['female', 'male'], [0, 1])
df['Embarked'] = df['Embarked'].replace(['S', 'C', 'Q'], [1, 2, 3])
train = df.query('train == 1')
test = df.query('train == 0')

In [34]: # Drop missing values from the train set.
train.dropna(axis=0, inplace=True)
labels = train[target].values
train.drop(['train', target, 'Pclass'], axis=1, inplace=True)
test.drop(['train', target, 'Pclass'], axis=1, inplace=True)

In [35]: #Test and train split
X_train, X_val, Y_train, Y_val = train_test_split(train, labels, test_size=0.2, random_state=1)

In [36]: classifier = GaussianNB()
classifier.fit(X_train, Y_train)

Out[36]: GaussianNB(priors=None, var_smoothing=1e-09)

In [37]: y_pred = classifier.predict(X_val)

# Summary of the predictions made by the classifier
print(classification_report(Y_val, y_pred))
print(confusion_matrix(Y_val, y_pred))
# Accuracy score
from sklearn.metrics import accuracy_score
print('accuracy is', accuracy_score(Y_val, y_pred))

precision    recall   f1-score   support
0.0         0.79      0.80      0.80       85
1.0         0.70      0.69      0.70       58

accuracy          0.76      143
macro avg       0.75      0.74      0.75      143
weighted avg     0.75      0.76      0.75      143

[[68 17]
 [18 40]]
accuracy is 0.7552447552447552
```

Question 3

(Glass Dataset)

1. Implement Naïve Bayes method using scikit-learn library.
 - a. Use the glass dataset available in Link also provided in your assignment.
 - b. Use train_test_split to create training and testing part.
2. Evaluate the model on testing part using score and classification_report(y_true, y_pred)
 1. Implement linear SVM method using scikit library
 - a. Use the glass dataset available in Link also provided in your assignment.
 - b. Use train_test_split to create training and testing part.
 2. Evaluate the model on testing part using score and

```
In [38]: glass=pd.read_csv("glass.csv")
glass.head()
```

```
Out[38]:
```

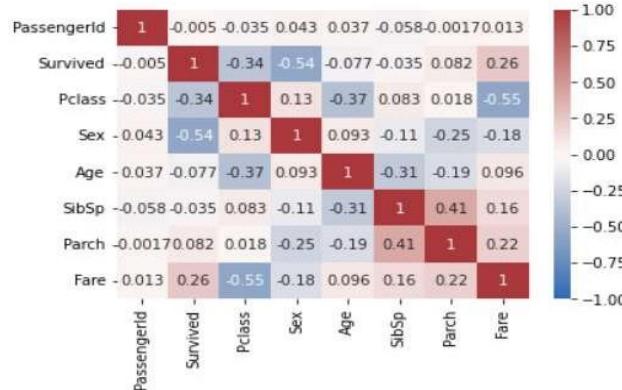
	RI	Na	Mg	Al	Si	K	Ca	Ba	Fe	Type
0	1.52101	13.64	4.49	1.10	71.78	0.06	8.75	0.0	0.0	1
1	1.51761	13.89	3.60	1.36	72.73	0.48	7.83	0.0	0.0	1
2	1.51618	13.53	3.55	1.54	72.99	0.39	7.78	0.0	0.0	1
3	1.51766	13.21	3.69	1.29	72.61	0.57	8.22	0.0	0.0	1
4	1.51742	13.27	3.62	1.24	73.08	0.55	8.07	0.0	0.0	1

```
In [39]: glass.corr().style.background_gradient(cmap="Reds")
```

```
Out[39]:
```

	RI	Na	Mg	Al	Si	K	Ca	Ba	Fe	Type
RI	1.000000	-0.191885	-0.122274	-0.407326	-0.542052	-0.289833	0.810403	-0.000386	0.143010	-0.164237
Na	-0.191885	1.000000	-0.273732	0.156794	-0.069809	-0.266087	-0.275442	0.326603	-0.241346	0.502898
Mg	-0.122274	-0.273732	1.000000	-0.481799	-0.165927	0.005396	-0.443750	-0.492262	0.083060	-0.744993
Al	-0.407326	0.156794	-0.481799	1.000000	-0.005524	0.325958	-0.259592	0.479404	-0.074402	0.598829
Si	-0.542052	-0.069809	-0.165927	-0.005524	1.000000	-0.193331	-0.208732	-0.102151	-0.094201	0.151565
K	-0.289833	-0.266087	0.005396	0.325958	-0.193331	1.000000	-0.317836	-0.042618	-0.007719	-0.010054
Ca	0.810403	-0.275442	-0.443750	-0.259592	-0.208732	-0.317836	1.000000	-0.112841	0.124968	0.000952
Ba	-0.000386	0.326603	-0.492262	0.479404	-0.102151	-0.042618	-0.112841	1.000000	-0.058692	0.575161
Fe	0.143010	-0.241346	0.083060	-0.074402	-0.094201	-0.007719	0.124968	-0.058692	1.000000	-0.188278
Type	-0.164237	0.502898	-0.744993	0.598829	0.151565	-0.010054	0.000952	0.575161	-0.188278	1.000000

```
In [40]: sns.heatmap(matrix, annot=True, vmax=1, vmin=-1, center=0, cmap='vlag')
plt.show()
```



```
In [41]: features = ['Rl', 'Na', 'Mg', 'Al', 'Si', 'K', 'Ca', 'Ba', 'Fe']
target = 'Type'
```

```
X_train, X_val, Y_train, Y_val = train_test_split(glass[:::-1], glass['Type'], test_size=0.2, random_state=1)
classifier = GaussianNB()
classifier.fit(X_train, Y_train)

y_pred = classifier.predict(X_val)

# Summary of the predictions made by the classifier
print(classification_report(Y_val, y_pred))
print(confusion_matrix(Y_val, y_pred))
# Accuracy score

print('accuracy is', accuracy_score(Y_val, y_pred))
```

	precision	recall	f1-score	support
1	0.90	0.95	0.92	19
2	0.92	0.92	0.92	12
3	1.00	0.50	0.67	6
5	0.00	0.00	0.00	1
6	1.00	1.00	1.00	1
7	0.75	0.75	0.75	4
accuracy			0.84	43
macro avg	0.76	0.69	0.71	43
weighted avg	0.89	0.84	0.85	43

```
[[18  1  0  0  0  0]
 [ 1 11  0  0  0  0]
 [ 1  0  3  2  0  0]
 [ 0  0  0  0  0  1]
 [ 0  0  0  0  1  0]
 [ 0  0  0  1  0  3]]
accuracy is 0.8372093023255814
```

```
In [43]: from sklearn.svm import SVC, LinearSVC
classifier = LinearSVC()
classifier.fit(X_train, Y_train)

y_pred = classifier.predict(X_val)

# Summary of the predictions made by the classifier
print(classification_report(Y_val, y_pred))
print(confusion_matrix(Y_val, y_pred))
# Accuracy score
from sklearn.metrics import accuracy_score
print('accuracy is',accuracy_score(Y_val, y_pred))
```

	precision	recall	f1-score	support
1	1.00	0.89	0.94	19
2	0.46	1.00	0.63	12
3	0.00	0.00	0.00	6
5	0.00	0.00	0.00	1
6	0.00	0.00	0.00	1
7	0.00	0.00	0.00	4
accuracy			0.67	43
macro avg	0.24	0.32	0.26	43
weighted avg	0.57	0.67	0.59	43

```
[[17  2  0  0  0]
 [ 0 12  0  0  0]
 [ 0  6  0  0  0]
 [ 0  1  0  0  0]
 [ 0  1  0  0  0]
 [ 0  4  0  0  0]]
accuracy is 0.6744186046511628
```

Justification:

We got better accuracy for Naïve Bayes method which is 0.8372093023255814. Naïve Bayes analysis works well with probabilistic concepts where as Linear SVM works better with linear regression logics. But to perform more accurately SVM requires large amounts of data to train and test the data. So, due to the amount of data Naïve Bayes algorith gives better accuracy compared to Linear SVM.